Reading the Order Flow: Information Extraction and Leakage in Learned Agents for the Trading Game Figgie

Nicholas Erup Larsen **Teammate 2** **Teammate 3**
Course 80245013, Tsinghua University
nicholaseruplarsen@gmail.com

Abstract

Figgie is a four-player card trading game in which the scoring ("goal") suit is hidden and must be inferred. All published Figgie agents, including the only reinforcement-learning attempt, infer the goal suit from the player's own hand alone, yet every trade made by an opponent also carries information about what that opponent holds. We propose to measure how much information the public order flow reveals, whether learned agents extract it, whether they leak their own private information through their actions, and whether leakage explains exploitability.

1 Problem and motivation

In Figgie, 40 cards are dealt from one of 12 equally likely decks; the goal suit is the same-colour partner of the 12-card suit and pays $10 per card plus a pot bonus to whoever holds the most. Players trade continuously, so a player who eagerly buys a suit is likely to hold, and believe in, that suit. Existing agents—figment's Bayesian market maker, Figgie-bot's exact inference, and the CS224R PPO agent [1]—compute or approximate the goal-suit posterior given only their own hand. We ask three questions:

1. **Revelation.** How much does the public trade history reveal about the goal suit beyond a single hand, and how quickly over a 240-tick game?
2. **Extraction.** Do agents trained with recurrent or attention-based policies learn to use that information, i.e. does their internal belief go beyond the exact hand-only posterior?
3. **Leakage and exploitability.** How much does each agent's own action stream reveal about its hand, and do the leakiest agents lose the most to a learned best response?

Elo and profit against a fixed field are the only metrics reported for Figgie agents, and they say nothing about *why* an agent wins or how robust it is. Information extraction from order flow and information leakage through one's own quotes are the central concerns of market microstructure; Figgie is a small, fully simulable market in which both can be measured against exact ground truth (the true deck and every hand). The study also exercises the core of the course: recurrent and transformer sequence models, representation learning via probes, and reinforcement learning with self-play.

2 Data

We use no external dataset. All data is generated by simulation in our engine, which implements the standard rules and the CS224R discretisation (240 ticks, 257 actions, prices 1–30). We have

Preprint. Under review.

already replicated four public sources on this engine: the CS224R baselines and PPO pipeline [1], figment [2], Figgie-bot [3] and figgie-auto [4]. The figment and Figgie-bot ports reproduce every decision of the original code (144,000/144,000 and 64,000/64,000 identical), and our PPO agent matches or exceeds every number in the CS224R report, including a 52% win rate (54% reported) against its strongest heuristic table. Each training or probing dataset is a set of simulated games with full logs: the deck, all hands, and every quote and trade.

3 Related work and relevance to the field

Market microstructure. Kyle [6] and Glosten–Milgrom [7] model how informed traders' orders reveal private information and how market makers protect themselves against it. Those models assume rational agents with known strategies. We measure revelation and leakage for learned and heuristic agents in a market where ground truth is exact, which is impossible on real exchanges.

Multi-agent learning. Exploitability is the standard robustness measure in two-player zero-sum games, but it has no closed form in multi-player general-sum games such as Figgie or six-player poker [8]. Learned best responses [9] are the practical substitute; we pair them with an explanation of *why* an agent is exploitable.

Interpretability of RL agents. Probing has shown that AlphaZero encodes human chess concepts [10], using linear probes [11]. Figgie offers an exact Bayesian target for the probe, so we can say not only whether a belief is encoded but how close it is to optimal.

4 Preliminary approach

4.1 Baselines

Belief: the exact hand-only Bayesian posterior over the 12 decks, which every published agent approximates. **Agents:** the CS224R PPO agent and the Bayesian market makers from figment and Figgie-bot, all replicated and verified on our engine; these are the reference points for strength, leakage and exploitability.

4.2 Method

Revelation bound. Train a supervised transformer that reads the public trade history up to tick t and predicts the goal suit. Compare its accuracy and log-loss over t with the exact hand-only posterior. The gap is an empirical upper bound on what any agent can gain from the order flow, measured separately against each population of opponents.

Extraction. Fit linear and small MLP probes from each learned agent's hidden state to (a) the exact hand-only posterior and (b) the true goal suit. Train new PPO agents whose observation includes the trade history, comparing a GRU, a transformer, and hand-crafted belief features, and track belief accuracy throughout training.

Leakage. For each agent, train a classifier that predicts that agent's hand composition from its actions alone; its accuracy above the prior is the agent's leakage.

Exploitability. For each agent, train a PPO best response at one seat with the other seats fixed, and estimate exploitability as the best response's profit gain over a no-information control. We decompose the gain into an information term (earlier correct goal-suit belief) and a pricing term, and test whether leakage predicts exploitability across agents. If time allows, we validate the estimator on small games where exact exploitability is computable (Kuhn and Leduc poker).

Expected results. (i) Goal-suit accuracy versus tick for hand-only, order-flow and probed agent beliefs; (ii) a table of Elo, profit, leakage and estimated exploitability for every agent; (iii) whether learned memory architectures close the gap to the revelation bound.

4.3 What we measure

Quantity	Definition
Revelation gap $\Delta_{\text{rev}}(t)$	Log-loss of the exact hand-only goal-suit posterior minus log-loss of the best predictor that also sees the public trade history up to tick t, in bits; reported for each opponent population.
Belief extraction $E_a(t)$	Log-loss of a probe on agent a's hidden state for the true goal suit, placed between the hand-only posterior ($E = 0$) and the revelation bound ($E = 1$).
Leakage L_a	Mutual information, in bits, between agent a's hand composition and its action stream, lower-bounded by the cross-entropy gain of a trained classifier over the prior.
Exploitability $\hat{\varepsilon}_a$	Profit per game of a learned best response at one seat against three copies of a, minus that of a no-information control. Decomposed into an information term and a pricing term.
Strength	Mean profit per game with 95% CIs, win rate and Elo over rotated mixed tables (already measured in the replication).

The central tests are: does $E_a > 0$ for any learned agent; does E_a increase with history-aware architectures (transformer $>$ GRU $>$ none); and is the rank correlation between L_a and $\hat{\varepsilon}_a$ positive across agents.

4.4 What is new

- Prior Figgie work ranks agents only by profit or Elo against a fixed field: hand-written strategies in a discrete-event simulation [5], Bayesian and evolutionary bots [2, 3], and one PPO agent [1]. None measures what the order flow reveals or what an agent leaks.
- No published Figgie agent conditions its belief on opponents' trades; all use the player's own hand. We test whether learning closes that gap.
- No exploitability estimate exists for any Figgie agent, and none for a multi-player trading game that we are aware of.
- We release a single engine with four replicated sources whose decisions are verified against the original code, so future work can compare against the same agents.

Timeline. Midterm (3 weeks): replication results, revelation bound, first probe of PPO. Final (8 further weeks): history-aware PPO variants, leakage measurement, best-response exploitability and the leakage–exploitability analysis.

References

[1] D. Li Yang. Reinforcement learning for Figgie: learning negotiation as a first-class skill. Stanford CS224R project report, 2026.

[2] mneha05. figment: a self-play market-making arena for Figgie-style card markets. `https://github.com/mneha05/figment`.

[3] revilo-25. Figgie-bot. `https://github.com/revilo-25/Figgie-bot`.

[4] 0xDub. figgie-auto. `https://github.com/0xDub/figgie-auto`.

[5] C. DiSilvio, A. Luo, and A. Ozerov. Traders in a strange land: agent-based discrete-event market simulation of the Figgie card game. arXiv:2110.00879, 2021.

[6] A. S. Kyle. Continuous auctions and insider trading. *Econometrica*, 53(6), 1985.

[7] L. R. Glosten and P. R. Milgrom. Bid, ask and transaction prices in a specialist market with heterogeneously informed traders. *J. Financial Economics*, 14(1), 1985.

[8] N. Brown and T. Sandholm. Superhuman AI for multiplayer poker. *Science*, 365(6456), 2019.

[9] F. Timbers et al. Approximate exploitability: learning a best response. In *IJCAI*, 2022.

[10] T. McGrath et al. Acquisition of chess knowledge in AlphaZero. *PNAS*, 119(47), 2022.

[11] G. Alain and Y. Bengio. Understanding intermediate layers using linear classifier probes. arXiv:1610.01644, 2016.